

A Statistical Approach to Used Car Price Prediction

12/09/2024

Steven Qie
*Statistics and Computer
Science*
*University of Illinois
Urbana-Champaign*
qie2@illinois.edu

Brian Gong
*Statistics and Computer
Science*
*University of Illinois
Urbana-Champaign*
brianhg2@illinois.edu

William Yeh
*Statistics and Computer
Science*
*University of Illinois
Urbana-Champaign*
wy16@illinois.edu

Introduction

With the used car market being significantly larger than the new car market, many consumers are realizing that used cars provide a more affordable option. It plays a significant role in the growth and stability of the U.S. economy, driven by changing consumer preferences, economic factors, and the availability of certain cars. Accurately predicting the price of a used car is a challenging but essential task for buyers, sellers, and market analysts/economists alike.

This report aims to develop various predictive models for used car prices using the Used Car Price Prediction Dataset from Kaggle. This dataset comprises of 4,009 data points, representing unique vehicle listings, as well as nine distinct features that serve as key indicators influencing the value of a used car. We follow a very structured and standard approach, including data exploration, preprocessing, model training, and evaluation using relevant performance metrics. By leveraging these methods, we aim to uncover valuable insights into the world of automobiles and the various factors that are driving used car prices.

Need a section on key findings. . . .

Abstract—

white space

white space

white space

white space

white space

white space

white space

white space

white space

We utilized AI tools in this report to enhance and assist in our writing. These tools helped play a big role in ensuring clarity, conciseness, and professionalism. We also utilized AI tools to help us with syntax help when writing code in R, as well as discovering potential bugs in our code.

Literature Review

This literature review aims to summarize key findings and approaches from a few noteworthy research papers focused on used car price prediction.

“Price Prediction of Used Cars Using Machine Learning”, written by Chuyang Jin of the University of Sydney, presents a model that can predict a used vehicle’s price given their year of production, mileage, tax, miles per gallon. He hopes that his model can benefit and save time for both sellers and buyers who are looking to sell or search for second-hand vehicles. Jin used a CSV dataset containing 100,000 records of used cars in the UK, focusing specifically on the Mercedes brand. The nine factors that he considered were the following: model, year, selling price, transmission, mileage, fuel type, tax, miles per gallon (mpg), and engine size. While doing exploratory data analysis and preprocessing, Jin noted that many many predictors had skewed distributions. For example, the overwhelming majority of prices fell in the 0-75,000 range, limiting the model’s potential effectiveness for higher price ranges. Jin deemed these data points as outliers and excluded them to ensure that the model would be more accurate and usable. After testing various forms of regression, namely linear, polynomial, SVR, Decision Trees, and Random Forests, Jin found Random Forest Regression yielded the best R squared value of 0.90416.

“Used Car Price Prediction using Machine Learning: A Case Study”, written by Mustapha Hankar, Marouane Birjali, and Abderrahim Beni-Hssane, applies several supervised machine learning algorithms to predict used car price prices based on features from a dataset collected from an online eCommerce website called Avito. During preprocessing, the authors of this paper performed recursive feature elimination to maintain only the most relevant features to car prices: year of manufacture, mileage, mark, fuel type, fiscal power, and model. Along with a baseline multiple linear regression model, the study also looked at K-nearest neighbors, Random Forest, Gradient Boosting, and Artificial Neural Networks. The study utilized 2 different performance metrics, R^2 and RMSE, and concluded that the Gradient Boosting Regression Model achieved the best results, with a R^2 of 0.8 and RMSE of 44516.20.

“Car Price Prediction using Supervised and Unsupervised Learning Models and Deep Learning” by Thomas Nsiah approached the problem of car price prediction from a supervised and unsupervised lenses. While supervised models allow a consumer to understand the key factors and predictors that influence pricing of used cars, unsupervised learning oftentimes uncovers hidden connections and patterns within the data. In his paper, Nsiah used a mock dataset of 50,000 UK second hand car sales with features similar to the previous 2 studies, such as model, engine size, fuel type, year, and mileage. Supervised learning models that Nsiah tried included simple linear regression, polynomial regression, and random forest, evaluated using mean absolute error (MAE) and R-squared metrics. He concluded that out of the supervised models, random forest performed best with an R-squared of 0.99849 and a MAE of 289.0691. For unsupervised learning techniques, Nsiah applied K-Means and DBSCAN clustering to identify price patterns, evaluated using the Davis Bouldin Index and the Silhouette Coefficient. He concluded that K-Means clustering for the year of manufacture vs price produced the best clustering results.

Overall, these three studies demonstrate the effectiveness that machine learning can have on accurately predicting used car prices. The next section will outline our own approach and findings.

Citations:

- C. Jin, “Price Prediction of Used Cars Using Machine Learning,” in 2021 IEEE International Conference on Emergency Science and Information Technology (ICESIT), Chongqing, China, 2021, pp. 223-230, doi: 10.1109/ICESIT53460.2021.9696839.
- M. Hankar, M. Birjali, and A. Beni-Hssane, “Used Car Price Prediction using Machine Learning: A Case Study,” in 2022 11th International Symposium on Signal, Image, Video and Communications (ISIVC), El Jadida, Morocco, 2022, pp. 1-4, doi: 10.1109/ISIVC54825.2022.9800719.
- T. Nsiah, “Car Price Prediction using Supervised and Unsupervised Learning Models and Deep Learning,” unpublished, 2024.

Data Processing and Summary Statistics

First, we will import the dataset and libraries into our workspace

#Preliminary Data Cleaning/Modifications First, we will removed the dollar sign and comma in price to enable numeric operations

Corrected the spelling of mileage from milage to mileage. Removed mi. and , to enable numeric operations.

The Engine columns contains very useful information such as the horsepower, displacement, cylinders, engine type, and fuel type. We turn these all into new columns.

```
##      brand      model model_year mileage fuel_type
## 1   Ford Utility Police Interceptor Base      2013    51000 Flex Fuel
## 2   Hyundai      Palisade SEL      2021    34742    <NA>
## 3   Lexus      RX 350 RX 350      2022    22372    <NA>
## 4 INFINITI      Q50 Hybrid Sport      2015    88900    Electric
## 5   Audi      Q3 45 S line Premium Plus      2021    9835    <NA>
## 6   Acura      ILX 2.4L      2016    136397    <NA>
##      transmission      ext_col int_col
## 1      6-Speed A/T      Black    Black
## 2 8-Speed Automatic      Moonlight Cloud    Gray
## 3      Automatic      Blue    Black
## 4      7-Speed A/T      Black    Black
## 5 8-Speed Automatic Glacier White Metallic    Black
## 6      F      Silver    Ebony.
##      accident clean_title price horsepower
## 1 At least 1 accident or damage reported      Yes 10300      300
## 2 At least 1 accident or damage reported      Yes 38005      NA
## 3      None reported      54598      NA
## 4      None reported      Yes 15500      354
## 5      None reported      34999      NA
## 6      None reported      14798      NA
##      displacement cylinders engine_type
## 1      3.7 6 Cylinder    <NA>
## 2      3.8    <NA>    DOHC
## 3      NA    <NA>    DOHC
## 4      3.5 6 Cylinder    <NA>
## 5      2.0    <NA>    DOHC
## 6      NA    <NA>    <NA>
```

Looking at each column's type and unique count

```
##      brand      model      model_year      mileage      fuel_type transmission
## "character" "character" "integer" "numeric" "character" "character"
##      ext_col      int_col      accident clean_title      price      horsepower
## "character" "character" "character" "character" "numeric" "numeric"
## displacement cylinders engine_type
## "numeric" "character" "character"

##      brand      model      model_year      mileage      fuel_type transmission
##      57      1898      NA      NA      7      62
##      ext_col      int_col      accident clean_title      price      horsepower
##      319      156      3      2      NA      NA
## displacement cylinders engine_type
##      NA      8      6
```

Let's examine columns that include NA or Empty String entries.


```
## [1] "fuel_type"      "accident"      "clean_title"   "horsepower"    "displacement"
## [6] "cylinders"      "engine_type"
```

#Analyzing categorical variables Categorical variables with various unique values include brand, model, transmission, ext_col, int_col. Let's examine all of them

First, we look at the "brand" and the "model" columns. Through analysis shown below, we have decided to omit both of these columns. Our reasoning and visualizations are shown below.

There are 57 unique brands with the frequency histogram not showing much dominance in a certain brand. To reduce the dimensionality, we will just omit this column

```
## [1] 57
```

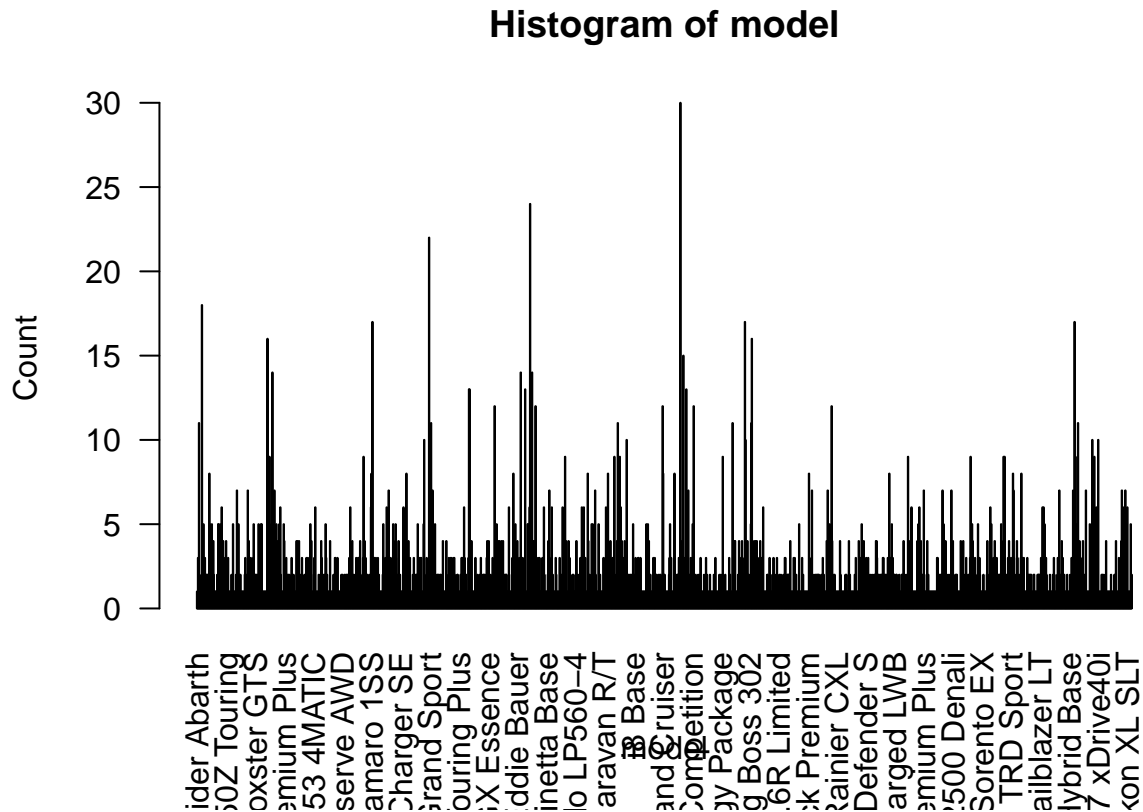


```
## # A tibble: 57 x 4
##   brand      medianprice averageprice count
##   <chr>          <dbl>         <dbl> <int>
## 1 Ford           32378.         36241.  386
## 2 BMW            32999         41072.  375
## 3 Mercedes-Benz  38598          52076.  315
## 4 Chevrolet      31992.         36723.  292
## 5 Porsche        59900          88751.  201
## 6 Audi           34498.         39907.  200
## 7 Toyota         27999          30026   199
```

```
## 8 Lexus          30000      35669.    163
## 9 Jeep           30000      31100.    143
## 10 Land          44924      55764.    130
## # i 47 more rows
```

This problem is seen even more in the model column. We also omit this column from the dataset


```
## [1] 1898
```



Now, let's examine colors. There are both intcol and extcol variables. Having too many unique color names can introduce noise into your classification model and make it harder for the model to generalize effectively. Grouping the colors into broader, more general categories can help improve model performance by reducing the dimensionality of the feature and making patterns more apparent.

```
## [1] 319
```

Histogram of ext_col




Let's apply the generalization function to simplify the different colors

```
## [1] "Black" "Other" "White" "Gray" "Gold" "Brown"
```

The same thing happens to int_col, but looking at the dataset we will have 4 categories.

```
## [1] 156
```

Histogram of interior color




```
## [1] "Black"      "Gray"      "Other"      "Beige/Ivory"
```

Examining the transmission column now

```
## [1] 62
```

Histogram of transmission



```
## # A tibble: 62 x 4
##   transmission      medianprice averageprice count
##   <chr>            <dbl>         <dbl> <int>
## 1 A/T              20500         31508.  1037
## 2 8-Speed A/T      39625         51126.   406
## 3 Transmission w/Dual Shift Mode 34000         54711.   398
## 4 6-Speed A/T      20900         25450.   362
## 5 6-Speed M/T      26450         39282.   248
## 6 Automatic       47541         63105.   237
## 7 7-Speed A/T      32999         47250.   209
## 8 8-Speed Automatic 41599         66072.   176
## 9 10-Speed A/T     57000         60915.   119
## 10 5-Speed A/T     15000         17607.    86
## # i 52 more rows
```

```
## [1] "A/T" "8-Speed A/T"
## [3] "Transmission w/Dual Shift Mode" "6-Speed A/T"
## [5] "6-Speed M/T" "Automatic"
## [7] "7-Speed A/T"
```

Now, we have looked at all the categorical variables with many many unique values, we will now one-hot encode the categorical variables. # One-hot encoding categorical variables

After looking at histograms for both Brand and Transmission, it seems Brand is more uniformly distributed while Transmission has a few salient categories. After exploring the categories of transmissions we found that the top 7 most frequent transmissions account for approximately 67-70% of the data points. Therefore we

will one hot encode these 7 categories + an “Other” category for Transmission for a total of 8 transmission categories. We will also one hot encode “fuel type” and “cylinders” since those are categorical variables as well.

Transmission

#Analyzing Null/Empty Values We will first look at the problem with NA and Empty values, something that this dataset has a lot of. We will first handle both NA and Empty “ ” values by replacing them to “NA” to make it easier to preprocess and analyze.

```
## [1] "fuel_type"      "accident"      "clean_title"   "horsepower"    "displacement"
## [6] "cylinders"      "engine_type"

## NULL

## NULL

##      [1] 300   NA 354 292 282 311 534 715 382 400 375 305 287 550 120
##     [16] 355 276 445 362 345 383 180 211 173 240 552 536 310 228 268
##     [31] 503 325 208 250 200 420 302 306 237 248 425 582 444 335 424
##     [46] 340 225 365 315 199 560 326 165 835 241 215 130 288 369 195
##     [61] 285 485 132 416 360 280 620 265 469 169 330 275 303 450 651
##     [76] 255 455 182 236 370 212 565 230 171 252 220 188 235 320 138
##     [91] 291 523 440 181 429 263 210 404 670 563 283 150 266 328 304
##    [106] 381 493 641 760 329 239 160 402 166 390 147 357 271 350 611
##    [121] 295 603 454 490 301 395 272 437 323 256 140 600 409 640 204
##    [136] 316 591 219 505 403 170 115 562 106 201 496 475 184 407 543
##    [151] 333 553 471 380 247 349 190 410 260 245 332 261 107 577 290
##    [166] 453 293 139 389 567 221 518 630 218 385 174 134 273 172 542
##    [181] 571 601 500 270 161 394 520 164 205 308 226 227 412 158 414
##    [196] 177 346 111 573 277 191 318 411 244 605 192 207 155 189 185
##    [211] 162 187 313 557 281 463 186 797 214 449 153 296 650 759 286
##    [226] 525 246 526 397 645 575 401 348 510 122 179 167 691 202 136
##    [241] 151 617 146 294 317 175 717 435 405 616 137 152 206 415 460
##    [256] 707 319 426 555 480 121 430 159 378 321 344 133 232 142 278
##    [271]  78 258 264 118  76 788 131 148 203 253 312 467 168 156 353
##    [286] 545 422 451 197 386 778 521 495 621 456 279 540 104 372 366
##    [301] 284 556 193 393 198 298 145 242 243  70 610 141 217 533 262
##    [316] 342 483 109 231 473 324 443 101 322 126 638 710 154 808 143
##    [331] 602 363 178 580 624 379 502 470 1020 572 702 660 341 222 729
##    [346] 417 482 224 176

##      [1] 3.70 3.80   NA 3.50 2.00 4.40 5.20 3.00 5.00 3.60 2.20 5.30 5.70 2.40 2.70
##     [16] 6.00 4.00 1.50 6.10 1.60 2.90 3.30 3.40 2.50 1.80 6.20 4.30 6.75 5.50 5.60
##     [31] 6.30 5.40 6.70 4.60 4.50 4.70 1.30 2.30 3.20 5.80 6.80 6.40 8.00 4.20 1.20
##     [46] 3.90 1.70 7.00 2.80 6.60 1.40 4.80 7.40 5.90 8.10 6.50 8.40 0.65 8.30 2.10
##     [61] 7.30 1.00

## NULL

## NULL
```

```
## [1] 810

## [1] 396

##
## 0.65    1  1.2  1.3  1.4  1.5  1.6  1.7  1.8    2  2.1  2.2  2.3  2.4  2.5  2.7
##      5    1    3    8   16   38   59    1   46  471    2    5   35   99  175   46
##  2.8  2.9    3  3.2  3.3  3.4  3.5  3.6  3.7  3.8  3.9    4  4.2  4.3  4.4  4.5
##      5   16  432   31   26   30  333  235   62  105   15  182   26   15   82    2
##  4.6  4.7  4.8    5  5.2  5.3  5.4  5.5  5.6  5.7  5.8  5.9    6  6.1  6.2  6.3
##     70   54   28  112   29  104   23   28   35  129    3    4   67    4  173    6
##  6.4  6.5  6.6  6.7  6.75  6.8    7  7.3  7.4    8  8.1  8.3  8.4
##     28    7   26   52    2    7    3    4    1    1    2    3    1


##      model_year      mileage      fuel_type      transmission
##      Min.      :1974      Min.      :   100      Length:4009      Length:4009
##      1st Qu.:2012      1st Qu.: 23044      Class :character      Class :character
##      Median :2017      Median : 52775      Mode  :character      Mode  :character
##      Mean   :2016      Mean   : 64718
##      3rd Qu.:2020      3rd Qu.: 94100
##      Max.   :2024      Max.   :405000
##
##      ext_col      int_col      accident      clean_title
##      Length:4009      Length:4009      Length:4009      Length:4009
##      Class :character      Class :character      Class :character      Class :character
##      Mode  :character      Mode  :character      Mode  :character      Mode  :character
##
##
##
##
##      price      horsepower      displacement      cylinders
##      Min.      : 2000      Min.      : 70.0      Min.      :0.650      Length:4009
##      1st Qu.: 17200      1st Qu.: 248.0      1st Qu.:2.500      Class :character
##      Median : 31000      Median : 310.0      Median :3.500      Mode  :character
##      Mean   : 44553      Mean   : 332.3      Mean   :3.711
##      3rd Qu.: 49990      3rd Qu.: 400.0      3rd Qu.:4.700
##      Max.   :2954083      Max.   :1020.0      Max.   :8.400
##
##                      NA's      :810      NA's      :396
##
##      engine_type
##      Length:4009
##      Class :character
##      Mode  :character
##
##
##
##
```

There are five columns with empty strings/NA values. Let's examine all five of them to discover if we can find any patterns.

horsepower

```
## [1] 348
```

```
## [1] 810
```



```
## [1] 0
```


Since there are 348 unique values in horsepower, we can consider horsepower as a continuous variable rather than categorical. However, there are 810 null values in a dataset with 4009 entries which is over 20% null values. This is too many to simply drop, so we want to perform some form of imputation. Looking at the distribution of horsepowers, we can see that the median is a good representative approximation for the distribution so we will use **median imputation**.

displacement (engine size)

```
## [1] 61
```

```
## [1] 396
```

Histogram of Displacement



```
## [1] 0
```

There are 61 unique values in displacement (engine size). Although these appear to be discretized measurements (ex: size = 0.8 or size = 3.71 may not make sense), we can treat it as a more continuous predictor for now. There are 396 null values in displacement which is just under 10% null values, so we could consider dropping these. However since the median already exists in the dataset (median = 3.5) we can also proceed with median imputation which is what we did.

```
## # A tibble: 7 x 4
##   fuel_type      medianprice averageprice count
##   <chr>          <dbl>         <dbl> <int>
## 1 Diesel        45450         48878.   114
## 2 Electric      42000.         46884.   238
## 3 Flex Fuel     18650         22156.   128
## 4 Gasoline      27000         38733.  2731
## 5 Hybrid        37999         45063.    17
## 6 Plug-In Electric/Gas 44945         45946.    34
## 7 <NA>          41599         68192.   747
```

The NA values for fuel_type have a higher median price and average price than other types, and makes up a significant count of observations so we are going to treat it as a separate category.

```
#cylinder
```

```
## # A tibble: 8 x 4
```

```
##   cylinders  medianprice averageprice count
##   <chr>         <dbl>         <dbl> <int>
## 1 10 Cylinder    100000         166530.    23
## 2 12 Cylinder     81330         140259.    37
## 3 3 Cylinder     32000         45281.    13
## 4 4 Cylinder     19000         22476.   739
## 5 5 Cylinder     10150.         18584.    20
## 6 6 Cylinder     27999         35935.  1225
## 7 8 Cylinder     34500         46401.  1007
## 8 <NA>          42599         64844.   945
```

#accident

```
## # A tibble: 3 x 4
##   accident                                medianprice averageprice count
##   <chr>                                <dbl>         <dbl> <int>
## 1 ""                                36500         50788.   113
## 2 "At least 1 accident or damage reported" 20900         28832.   986
## 3 "None reported"                    35668.         49638.  2910
```

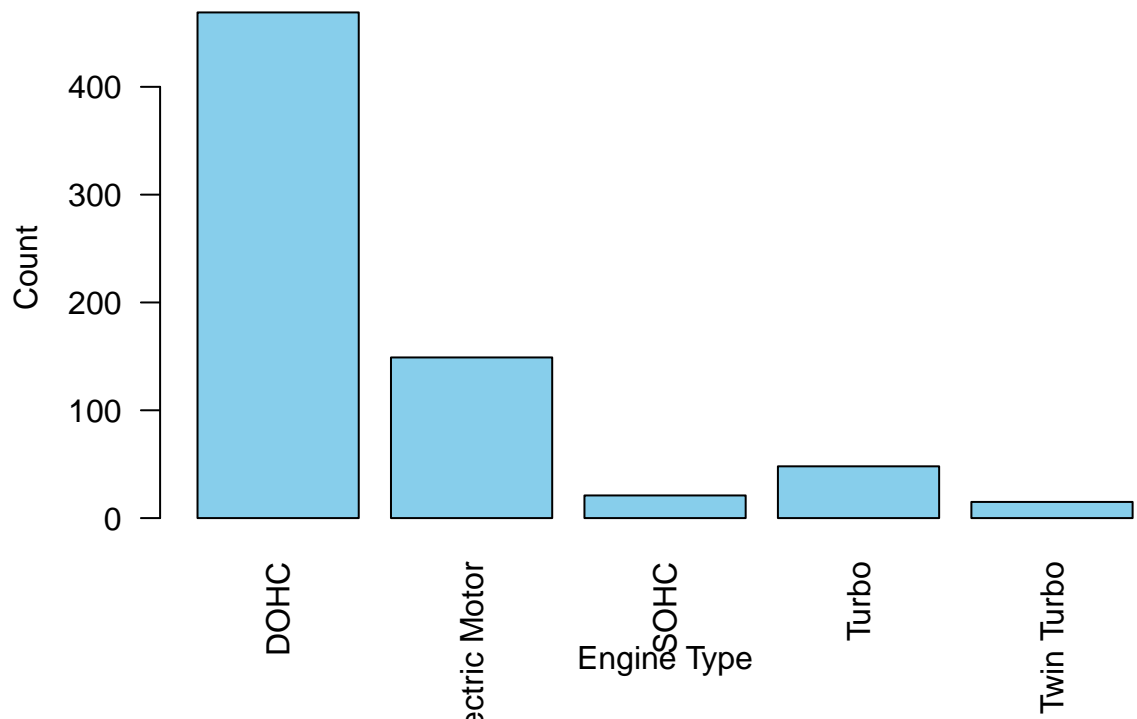
The NA/Empty values for accident exhibit very similar properties to the None reported category, with median price and average price being pretty similar, not to mention a very small percentage of data is represented by this value. Therefore, we replace and combine these observations with the None reported category. Because accident only has 2 unique values now, no accidents and 1 or more accidents, we changed it to 1,0 to be useful for models.

```
## # A tibble: 2 x 4
##   clean_title medianprice averageprice count
##   <chr>         <dbl>         <dbl> <int>
## 1 ""          42996.         60695.   596
## 2 "Yes"       29000         41734.  3413
```

The NA values for clean_title clearly have a significantly higher median price and will be treated as a separate category. We apply similar reasoning from accident to clean_title. Since there is only “Yes” and NA, we treat all the yes’s to 1 and all the NA values to 0.

```
## [1] 1 0
```

Histogram of engine_type



#engine type

```
## # A tibble: 6 x 4
##   engine_type medianprice averageprice count
##   <chr>         <dbl>         <dbl> <int>
## 1 DOHC          39244          77951.  469
## 2 Electric Motor 47800          54439.  149
## 3 SOHC          38998          38676.   21
## 4 Turbo         49940.          51767.   48
## 5 Twin Turbo     85998          89258.   15
## 6 <NA>          28250          39101.  3307
```

#Removing Outliers We remove outliers with 1.5*IQR value.

```
## [1] "Number of outliers: 244 and average price of these cars: 214826.76"
```

```
##   model_year mileage fuel_type transmission
##   Min. :1992 Min. : 100 Length:3765 Length:3765
##   1st Qu.:2012 1st Qu.: 26600 Class :character Class :character
##   Median :2017 Median : 57237 Mode :character Mode :character
##   Mean :2015 Mean : 68075
##   3rd Qu.:2020 3rd Qu.: 97000
##   Max. :2024 Max. :405000
##   ext_col int_col accident clean_title
##   Length:3765 Length:3765 Min. :0.0000 Min. :0.0000
##   Class :character Class :character 1st Qu.:0.0000 1st Qu.:1.0000
##   Mode :character Mode :character Median :0.0000 Median :1.0000
```

```
##                               Mean    :0.2595   Mean    :0.8608
##                               3rd Qu.:1.0000   3rd Qu.:1.0000
##                               Max.     :1.0000   Max.     :1.0000
##      price      horsepower      displacement      cylinders
## Min.    : 2000   Min.      : 70.0   Min.     :0.650   Length:3765
## 1st Qu.:16500   1st Qu.: 263.0   1st Qu.:2.500   Class :character
## Median :29600   Median : 310.0   Median :3.500   Mode  :character
## Mean    :33518   Mean     : 320.9   Mean     :3.648
## 3rd Qu.:45500   3rd Qu.: 375.0   3rd Qu.:4.400
## Max.    :99000   Max.     :1020.0   Max.     :8.300
## engine_type
## Length:3765
## Class :character
## Mode  :character
##
##
##
```

#turning each categorical column into a factor type

#one hot encoding Some models will require one hot encoding. For these models, we create a new dataset and apply this one hot encoding

#Final Summary Statistics

```
## [1] 3765    46
```

```
##      model_year      mileage      fuel_type.Diesel fuel_type.Electric
## Min.    :1992   Min.      : 100   Min.     :0.00000   Min.     :0.00000
## 1st Qu.:2012   1st Qu.: 26600   1st Qu.:0.00000   1st Qu.:0.00000
## Median :2017   Median : 57237   Median :0.00000   Median :0.00000
## Mean    :2015   Mean     : 68075   Mean     :0.02895   Mean     :0.06135
## 3rd Qu.:2020   3rd Qu.: 97000   3rd Qu.:0.00000   3rd Qu.:0.00000
## Max.    :2024   Max.     :405000   Max.     :1.00000   Max.     :1.00000
## fuel_type.Flex Fuel fuel_type.Gasoline fuel_type.Hybrid fuel_type.NA
## Min.    :0.000   Min.     :0.0000   Min.     :0.00000   Min.     :0.0000
## 1st Qu.:0.000   1st Qu.:0.0000   1st Qu.:0.00000   1st Qu.:0.0000
## Median :0.000   Median :1.0000   Median :0.00000   Median :0.0000
## Mean    :0.034   Mean     :0.6887   Mean     :0.00425   Mean     :0.1737
## 3rd Qu.:0.000   3rd Qu.:1.0000   3rd Qu.:0.00000   3rd Qu.:0.0000
## Max.    :1.000   Max.     :1.0000   Max.     :1.00000   Max.     :1.0000
## fuel_type.Plug-In Electric/Gas transmission.6-Speed A/T
## Min.    :0.00000   Min.     :0.00000
## 1st Qu.:0.00000   1st Qu.:0.00000
## Median :0.00000   Median :0.00000
## Mean    :0.00903   Mean     :0.09535
## 3rd Qu.:0.00000   3rd Qu.:0.00000
## Max.    :1.00000   Max.     :1.00000
## transmission.6-Speed M/T transmission.7-Speed A/T transmission.8-Speed A/T
## Min.    :0.00000   Min.     :0.00000   Min.     :0.00000
## 1st Qu.:0.00000   1st Qu.:0.00000   1st Qu.:0.00000
## Median :0.00000   Median :0.00000   Median :0.00000
## Mean    :0.06348   Mean     :0.05206   Mean     :0.09854
## 3rd Qu.:0.00000   3rd Qu.:0.00000   3rd Qu.:0.00000
```

```

## Max. :1.00000 Max. :1.00000 Max. :1.00000
## transmission.A/T transmission.Automatic transmission.Other
## Min. :0.0000 Min. :0.00000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.00000 1st Qu.:0.0000
## Median :0.0000 Median :0.00000 Median :0.0000
## Mean :0.2667 Mean :0.05657 Mean :0.2709
## 3rd Qu.:1.0000 3rd Qu.:0.00000 3rd Qu.:1.0000
## Max. :1.0000 Max. :1.00000 Max. :1.0000
## transmission.Transmission w/Dual Shift Mode ext_col.Black ext_col.Brown
## Min. :0.00000 Min. :0.000 Min. :0.00000
## 1st Qu.:0.00000 1st Qu.:0.000 1st Qu.:0.00000
## Median :0.00000 Median :0.000 Median :0.00000
## Mean :0.09641 Mean :0.255 Mean :0.02125
## 3rd Qu.:0.00000 3rd Qu.:1.000 3rd Qu.:0.00000
## Max. :1.00000 Max. :1.000 Max. :1.00000
## ext_col.Gold ext_col.Gray ext_col.Other ext_col.White
## Min. :0.00000 Min. :0.0000 Min. :0.0000 Min. :0.000
## 1st Qu.:0.00000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.000
## Median :0.00000 Median :0.0000 Median :0.0000 Median :0.000
## Mean :0.01116 Mean :0.2653 Mean :0.2133 Mean :0.234
## 3rd Qu.:0.00000 3rd Qu.:1.0000 3rd Qu.:0.0000 3rd Qu.:0.000
## Max. :1.00000 Max. :1.0000 Max. :1.0000 Max. :1.000
## int_col.Beige/Ivory int_col.Black int_col.Gray int_col.Other
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.0000 Median :1.0000 Median :0.0000 Median :0.0000
## Mean :0.1392 Mean :0.5214 Mean :0.1246 Mean :0.2149
## 3rd Qu.:0.0000 3rd Qu.:1.0000 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max. :1.0000 Max. :1.0000 Max. :1.0000 Max. :1.0000
## accident clean_title price horsepower
## Min. :0.0000 Min. :0.0000 Min. : 2000 Min. : 70.0
## 1st Qu.:0.0000 1st Qu.:1.0000 1st Qu.:16500 1st Qu.: 263.0
## Median :0.0000 Median :1.0000 Median :29600 Median : 310.0
## Mean :0.2595 Mean :0.8608 Mean :33518 Mean : 320.9
## 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:45500 3rd Qu.: 375.0
## Max. :1.0000 Max. :1.0000 Max. :99000 Max. :1020.0
## displacement cylinders.10 Cylinder cylinders.12 Cylinder
## Min. :0.650 Min. :0.000000 Min. :0.000000
## 1st Qu.:2.500 1st Qu.:0.000000 1st Qu.:0.000000
## Median :3.500 Median :0.000000 Median :0.000000
## Mean :3.648 Mean :0.002922 Mean :0.005578
## 3rd Qu.:4.400 3rd Qu.:0.000000 3rd Qu.:0.000000
## Max. :8.300 Max. :1.000000 Max. :1.000000
## cylinders.3 Cylinder cylinders.4 Cylinder cylinders.5 Cylinder
## Min. :0.000000 Min. :0.0000 Min. :0.000000
## 1st Qu.:0.000000 1st Qu.:0.0000 1st Qu.:0.000000
## Median :0.000000 Median :0.0000 Median :0.000000
## Mean :0.003453 Mean :0.1958 Mean :0.005312
## 3rd Qu.:0.000000 3rd Qu.:0.0000 3rd Qu.:0.000000
## Max. :1.000000 Max. :1.0000 Max. :1.000000
## cylinders.6 Cylinder cylinders.8 Cylinder cylinders.NA engine_type.DOHC
## Min. :0.0000 Min. :0.000 Min. :0.0000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.0000 Median :0.000 Median :0.0000 Median :0.0000

```



```

## Mean      :0.3118      Mean      :0.251      Mean      :0.2242      Mean      :0.1039
## 3rd Qu.   :1.0000      3rd Qu. :1.000      3rd Qu. :0.0000      3rd Qu. :0.0000
## Max.      :1.0000      Max.      :1.000      Max.      :1.0000      Max.      :1.0000
## engine_type.Electric Motor engine_type.NA engine_type.SOHC
## Min.      :0.00000      Min.      :0.0000      Min.      :0.000000
## 1st Qu.   :0.00000      1st Qu. :1.0000      1st Qu. :0.000000
## Median    :0.00000      Median   :1.0000      Median   :0.000000
## Mean      :0.03825      Mean      :0.8369      Mean      :0.005578
## 3rd Qu.   :0.00000      3rd Qu. :1.0000      3rd Qu. :0.000000
## Max.      :1.00000      Max.      :1.0000      Max.      :1.000000
## engine_type.Turbo engine_type.Twin Turbo
## Min.      :0.00000      Min.      :0.000000
## 1st Qu.   :0.00000      1st Qu. :0.000000
## Median    :0.00000      Median   :0.000000
## Mean      :0.01275      Mean      :0.002656
## 3rd Qu.   :0.00000      3rd Qu. :0.000000
## Max.      :1.00000      Max.      :1.000000

```


Unsupervised Learning

Apply at least three clustering algorithms to the processed dataset. Determine the appropriate number of clusters and discuss the interpretability of these clusters. Do they hold any meaningful distinctions? Examine whether the clustering results are associated with your outcome variable.


1. KMeans Clustering

We decided to use kmeans to examine the relation between model_year and price, as we noticed a similar examination in one of the papers while doing the literature review. Because K-means utilizes distance metrics, we scale the data before clustering.

Silhouette Method for Optimal k



We decided to use the Silhouette Method to determine the optimal number of clusters. This method essentially uses distance measures calculating how close clusters are to themselves and how far away they are to other clusters to judge the optimal number of clusters. In this case, 2 has the highest average silhouette score so we will use $k=2$.

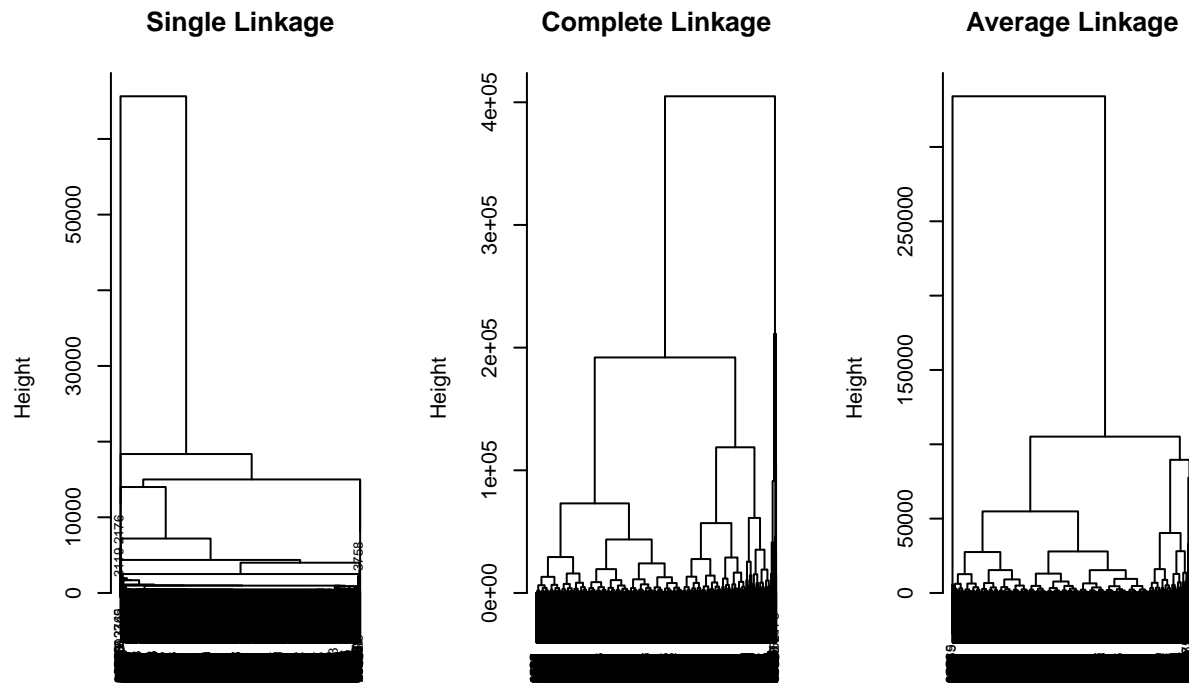


There seems to be a pretty solid relationship between a more recent model_year and higher price. Although the 2 clusters seem to be mostly dominated by model year, it's clear that the average price of cluster 2 is higher than cluster 1.

2. Hierarchical Clustering

Next, we will try hierarchical clustering with three different linkage methods(single, complete, and average) using euclidean distance. Hierarchical Clustering begins with each data point starting as its own cluster. The goal is to progressively group them together until there is only one group. The process involves choosing the closest two groups, calculated through a specific distance metric.

Removing non-numeric features as clustering requires numeric features. Also, removed the target feature price.



```
##
##      1      2
## 3674    91

## # A tibble: 2 x 7
##   cluster avg_price avg_model_year avg_accident avg_mileage avg_horsepower count
##   <fct>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl> <int>
## 1 1          34086.        2015.        0.254      64111.        322.  3674
## 2 2          10588.        2007.        0.484      228100.       267.   91
```


There are a lot of correlations here that make sense between the 2 clusters. Cluster 1, with a more recent avg_model_year, also has a lower avg_mileage and a lower avg_accident rate, probably because the car has been driven for less time, this cluster also has a much higher avg_price in comparison to cluster 2. The data isn't distributed very well however as a vast majority of the points sit in cluster 1, perhaps suggesting that hierarchical clustering isn't suitable for this dataset.

3. Spectral Clustering

Finally, we will try spectral clustering, which aims to group observations based on their proximity information. This method involves 2 main steps, the first being using the eigenvalues of a similarity matrix to perform dimension reduction, followed by applying a clustering algorithm like K-means.

```
## # A tibble: 2 x 6
##   cluster avg_model_year avg_mileage avg_accident avg_horsepower count
```

##	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>
## 1	1	2016.	60899.	0.246	325.	859
## 2	2	2016.	60429.	0.220	329.	141



Similar to Cluster 1, with a more recent `avg_model_year`, also has a lower `avg_mileage` and a lower `avg_accident` rate, this cluster also has a much higher `avg_price` in comparison to cluster 2. The distribution of data points between the 2 clusters seem to be more even in comparison to heirarchical clustering, meaning that perhaps spectral clustering is more suitable for this dataset.

Prediction Models

For all the supervised models below, we will split the data into training sets for model training and testing sets to evaluate performance and accuracy

1. Linear Model. There are mainly three possible linear models: Lasso, Ridge, and Elastic Net. We will try all three models and see which one performs the best. Lasso, Ridge, and Elastic Net all benefit from feature scaling because these models involve regularization. which will penalize the size of coefficients of the model to avoid overfitting. All 3 models also involving a tuning parameter, and so we will use k-fold cross validation to find the best parameters. `cv.glmnet` will automatically scale and center the data as well.

Training our ridge model

```
## [1] 1368.33
```

```
## [1] 12017.06
```

Training our lasso model

```
## [1] 82.02921
```

```
## [1] 11992.87
```

Training our elastic net model

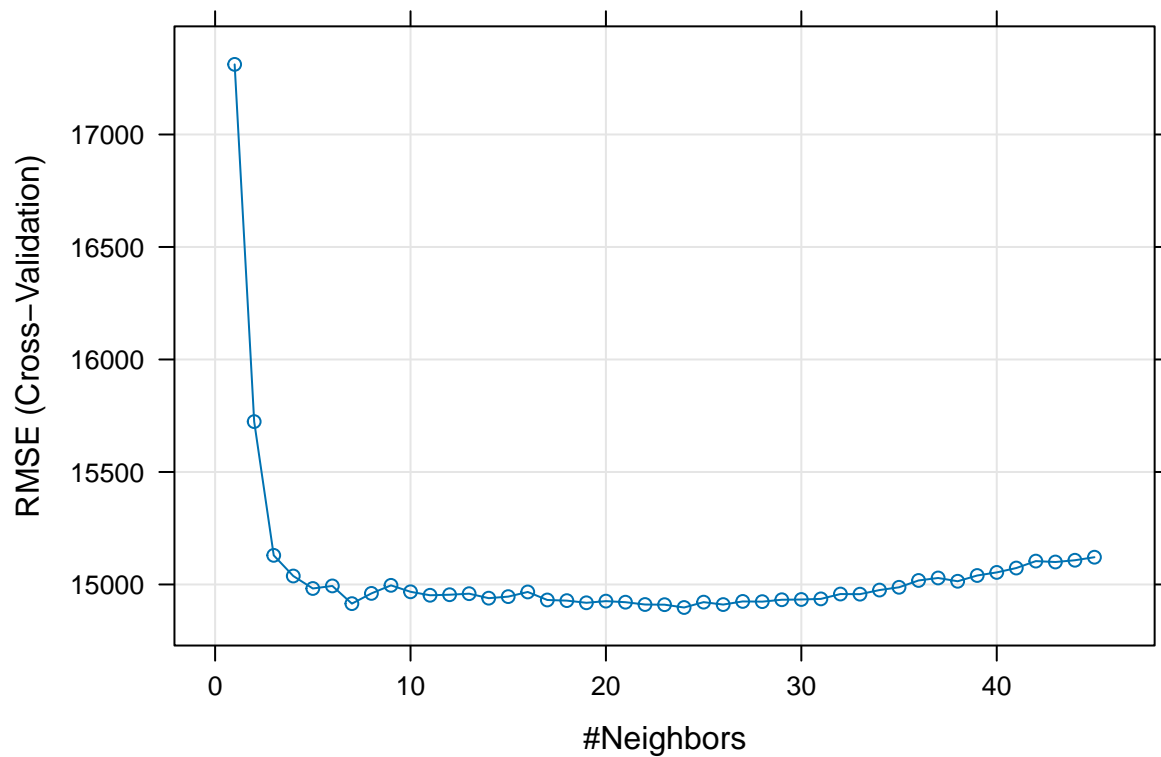
```
## [1] 149.4839
```

```
## [1] 11993.98
```

Out of our 3 linear models, Ridge performed the best, with a RMSE of 12261.19

2. K Nearest Neighbors(KNN) regression works by calculating the k nearest training set data points to the test point and predicting the target value by taking the average of their target values. KNN is sensitive to feature scaling, so we will need to scale the data. The reason behind this is for example, if one feature has ranges from 1-10 and another one has 1-10000, distance calculations will be biased and results will suffer as a result. KNN is also sensitive to the choice of k. To find the optimal value of k, we will perform k-fold cross validation.

```
## Number of components to retain: 32
```



```
## [1] "The best value of k based on cross-validation is: 24"
```

```
## [1] "Prediction error: 39760.6901508748"
```

3. Random Forest
4. SVM? does this count as a linear model
5. Gradient Boosting Regressor

Open-Ended Question/Conclusion

A researcher is interested in estimating the original price of the cars in our dataset as if they were brand new. To solve this problem, we built a machine learning model using features that we believed are most relevant to understanding depreciation. We utilized the following key features to model depreciation:

1. Brand: Different brands might exhibit varying depreciation patterns. For example, economy brands like Toyota may have a more linear depreciation curve, while luxury brands like BMW may see a steeper initial depreciation. Whether a brand is considered luxury or economy is useful as well, because luxury brands typically have higher starting prices than economy brands. Because of these reasons, we decided to include this feature even though it was omitted in our earlier prediction models.
2. Age: Cars will generally lose value over time as newer models with newer designs and updated features are released. Since we are only given the model_year in the dataset, which represents when the car was made, we created a new column called "Age" that is simply current_year - model_year + 1, which represents the number of years since the car was new.
3. Mileage: Generally, the higher the mileage, the lower the car's value. Increased miles often indicates more wear and tear on car components. High mileage on a car can also signal that more maintenance will be required to keep the car running smoothly, making it less attractive to potential buyers.
4. Accident History: This feature helps capture the impact of damage on a car's value because accidents are usually correlated with the integrity, safety, and reliability of the car. Many buyers will perceive cars with accident histories as less favorable compared to cars without accident histories.
5. Clean Title: Cars without a clean title may depreciate faster due to the higher perceived risks and uncertainties associated with its condition.

For practicality reasons, and to manage potential model complexity issues that arises from having to one-hot encode every single unique brand, we will train our model on a subset of the data, specifically cars belonging to the seven most common brands in our dataset: Ford, BMW, Mercedes-Benz, Chevrolet, Porsche, Audi, and Toyota. This subset also provides a good balance between luxury and economy brands.

Estimating brand-new car prices using only a dataset of used cars presents a couple of challenges. Because the model has never actually seen cars with zero age and zero mileage, it must infer/extrapolate what the original prices might have been, which can lead to potential inaccuracies. Additionally, external factors that the data set doesn't capture, such as inflation, shortages, and market competition, also influence the car's original price. For example, the first six weeks of 2022 saw a car shortage caused by factors such as the pandemic and a global semiconductor shortage, none of which are accounted for in the model.

Nevertheless, we aim to leverage these features to develop a machine learning model that can reasonably estimate the original prices of cars.

##	price	brand.Audi	brand.BMW	brand.Chevrolet
##	"numeric"	"numeric"	"numeric"	"numeric"
##	brand.Ford	brand.Mercedes-Benz	brand.Porsche	brand.Toyota
##	"numeric"	"numeric"	"numeric"	"numeric"
##	mileage	accident	clean_title	Age
##	"numeric"	"numeric"	"numeric"	"numeric"

lasso model

Original Price Estimation and Comparison

We selected one car each from Toyota, Ford, and Porsche and using the trained model, predicted the “brand-new” price for each car. Then, we searched online for the actual original release prices of these cars and compared them with the model’s predictions. R Code is hidden to make the report more easy to read.

```
## Randomly picked Toyota Car: 4Runner TRD Off Road. Year of Model: 2019
## Predicted Price: 66952.068
```

The original MSRP was approximately 40,395. Our model predicted the car to be around 26,000 higher than the actual price.

```
## Randomly picked Ford Car: Ford Maverick XLT. Year of Model: 2022
## Predicted Price: 67774.82
```

The original MSRP was approximately 29,541. Our model predicted the car to be around 37,000 higher than the actual price.

```
## Randomly picked Porsche Car: 718 Spyder Base. Year of Model: 2022
## Predicted Price: 111831.508
```

The original MSRP was approximately 98300. Our model predicted the car to be around 13500 higher than the actual price.